

MAHESWARA REDDY VARRA

DATA ENGINEER | Python • SQL • PySpark • AWS • Airflow

989-933-6619 | maheswara7569@gmail.com

linkedin.com/in/maheswara-reddy-varra | github.com/maheswara11

SUMMARY

Data engineering candidate with an M.S. in Information Systems and experience preparing, validating, and reporting business data. Built AWS lakehouse pipelines using Python, SQL, PySpark, and Airflow, with incremental ingestion, data quality checks, monitoring, and CI/CD. Developed a Data Quality Copilot for profiling and reviewing corrections to uploaded datasets.

TECHNICAL SKILLS

Languages and databases: Python, SQL, Java, PostgreSQL, MySQL

Data engineering: Apache Spark, PySpark, Databricks, dbt (SQL models, incremental models, snapshots), ETL/ELT, Apache Airflow, Pandas, data validation, dimensional modeling, star schema

AWS: S3, Glue, Glue Data Catalog, Athena, Lambda, IAM, CloudWatch, SNS, Lake Formation

Cloud platforms: AWS (project experience), Microsoft Azure (fundamentals, AZ-900), Google Cloud Platform (GCP fundamentals)

Tools: Git, GitHub Actions, CI/CD, Docker, Terraform, Linux, Tableau, Excel

PROFESSIONAL EXPERIENCE

Data Analyst | PMR Networks | India | March 2023 – August 2024

- Cleaned and transformed data with SQL and Excel to prepare consistent datasets for business reporting.
- Validated and reconciled records, investigated data discrepancies, and worked with stakeholders to resolve reporting issues.
- Built Tableau dashboards and reporting datasets to help stakeholders review business performance.

DATA ENGINEERING PROJECTS

Retail Sales Lakehouse | AWS S3, Glue, PySpark, Airflow, Athena, Terraform

- Built a Bronze–Silver–Gold lakehouse on Amazon S3, using AWS Glue and PySpark to transform raw retail data into curated datasets for Athena queries.
- Implemented watermark-based incremental processing and Airflow orchestration with retries and failure handling; registered datasets in the Glue Data Catalog.
- Configured CloudWatch monitoring and SNS alerts; used Terraform and GitHub Actions with OIDC for infrastructure and deployment workflows.

AI Data Quality Copilot | Python, Pandas, Streamlit, Ollama, Gen AI

- Developed CSV and JSON profiling to identify null values, duplicate primary keys, whitespace, and invalid email and date values, with failed-row details.
- Added user-approved cleaning on a working copy and before-and-after validation so users can review changes while preserving original data.
- Connected quality reports to a local LLM for explanations and provided cleaned-data downloads through a Streamlit interface.

AWS Serverless Data Processing Pipeline | Python, Amazon S3, AWS Lambda

- Built an S3-triggered Lambda workflow to process uploaded files, select title and rating columns, and write results to a separate processed-data location.

EDUCATION AND CERTIFICATION

M.S. in Information Systems | Central Michigan University | May 2026

GPA: 3.74/4.0 | Coursework: Database Management, Cloud Computing, Business Intelligence

B.E. in Electronics and Communication Engineering

Microsoft Certified: Azure Fundamentals (AZ-900)